

ReCell Project

Supervised Learning - Linear Regression

Yair Brama – August 2024

Contents / Agenda

- Executive Summary
- Business Problem Overview and Solution Approach
- EDA Results
- Data Preprocessing
- Model Performance Summary
- Appendix

Executive Summary - Business Context

The rising potential of this comparatively under-the-radar market of refurbished and used cellular devices fuels the need for an ML-based solution to develop a dynamic pricing strategy. ReCell, a startup aiming to tap the potential in this market, has hired me as a data scientist. In this project we will analyze the data provided and build a linear regression model to predict the price of a used phone/tablet and identify factors that significantly influence it. (See the data structure in the appendix part)

Performing Supervised Machine Learning – Linear Regression

The primary purpose of linear regression is to predict the value of the dependent variable (price of used devices) based on the values of the independent variables.

As it will show in the next slides, the model that is based on the given data can predict the price with an accuracy of ~81%, which meets our expectations for a reliable model

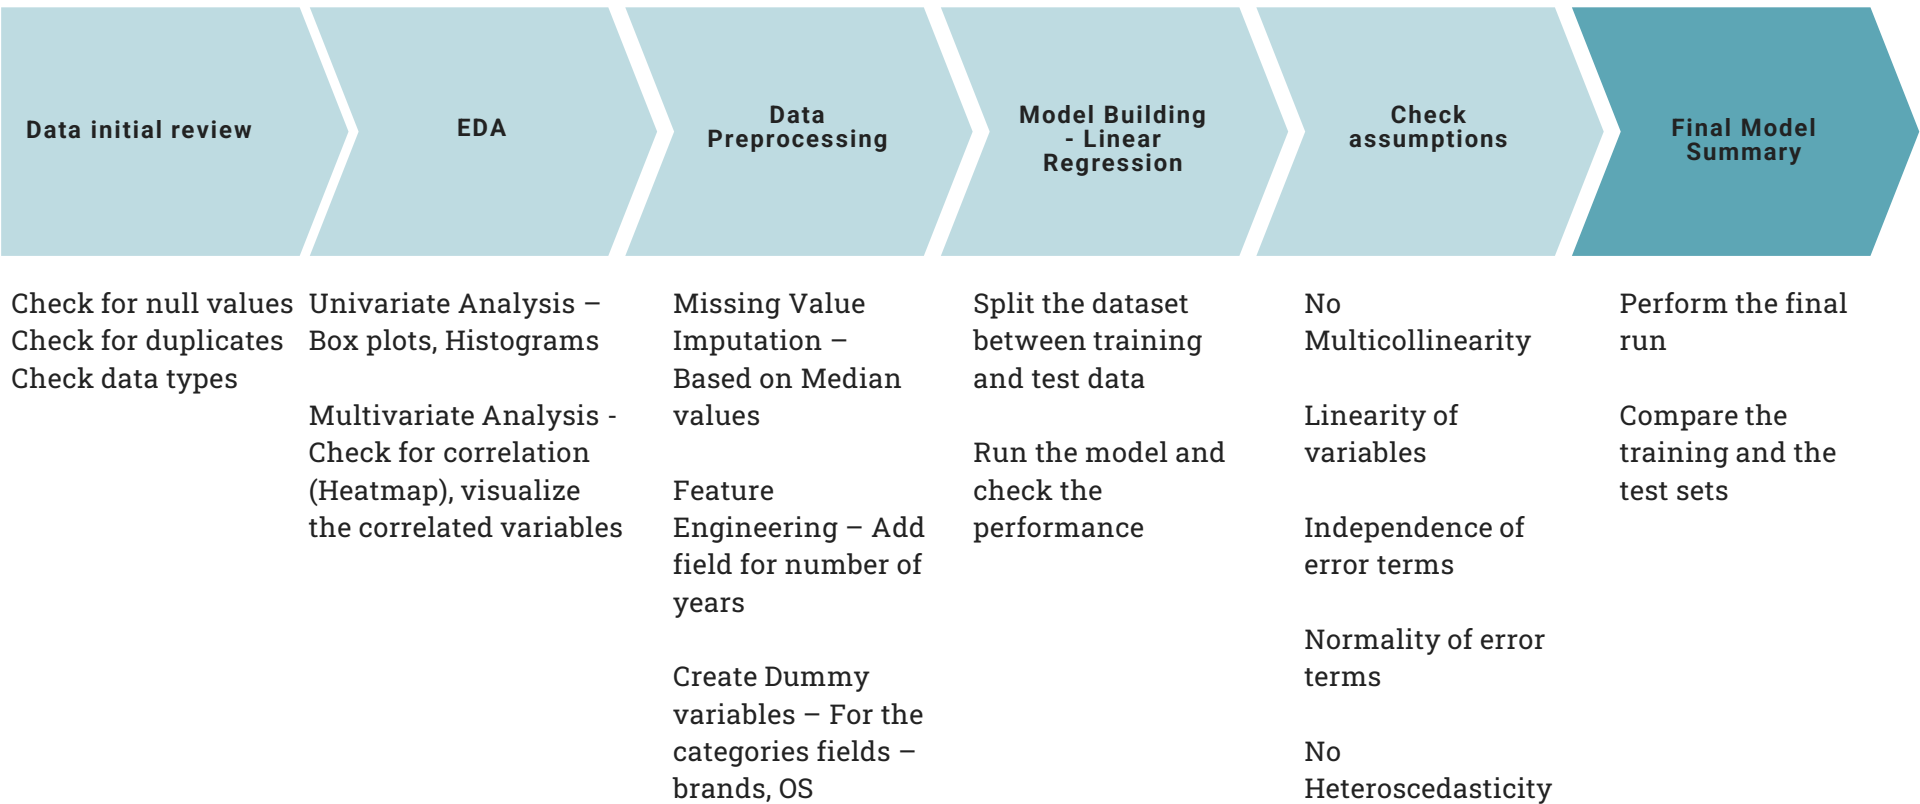
Executive Summary – Model Results

| Metric | Train Value | Test Value | Description |
|--------------------|-----------------|-----------------|----------------------------------------------------------------------------------------------------------------|
| R-squared | 0.816984 | 0.822131 | The training model is NOT underfitting, ~82% of the results can be predicted by the model |
| Adjusted R-squared | 0.81584 | 0.819518 | This value shows how well the model fits the data, considering the amount of variables |
| MAE | 0.194005 | 0.196185 | MAE indicates that our current model can predict the price within a mean error of ~0.2 units on the test data. |
| MAPE | 0.046678 | 0.047941 | MAPE of 4.8% shows we can predict the price of the test data within this range |
| RMSE | 0.249705 | 0.253286 | The train and test RMSE (as well as MAE) are comparable, so the model is not overfitting |

Equation of linear regression

normalized price of used device = $1.3946020282806701 + 0.01467795180139142 * (\text{main_camera_mp}) + 0.01171334063789519 * (\text{selfie_camera_mp}) + 0.018694083226400683 * (\text{ram}) + 9.434404942690962e-05 * (\text{battery}) + 0.4756460141277346 * (\text{normalized_new_price}) + -0.021515391520670724 * (\text{years_since_release}) + 0.05985704728434499 * (\text{brand_name_Alcatel}) + 0.14081623761562923 * (\text{brand_name_Karbonn}) + 0.0645434807324212 * (\text{brand_name_Lenovo}) + 0.13940662078054894 * (\text{brand_name_Microsoft}) + 0.07148832358895618 * (\text{brand_name_Nokia}) + 0.07361505473160784 * (\text{brand_name_Xiaomi}) + -0.1454720059578327 * (\text{os_Others}) + -0.0763163720532658 * (\text{5g_yes})$

Solution Approach



EDA Results - Data Overview

The data includes 3454 rows and 15 columns

There are 34 unique brands (e.g., Nokia, Apple), and 4 OS type (e.g., Android, iOS)

Data types:

| # | Column | Non-Null Count | Dtype |
|----|-----------------------|----------------|---------|
| 0 | brand_name | 3454 non-null | object |
| 1 | os | 3454 non-null | object |
| 2 | screen_size | 3454 non-null | float64 |
| 3 | 4g | 3454 non-null | object |
| 4 | 5g | 3454 non-null | object |
| 5 | main_camera_mp | 3275 non-null | float64 |
| 6 | selfie_camera_mp | 3452 non-null | float64 |
| 7 | int_memory | 3450 non-null | float64 |
| 8 | ram | 3450 non-null | float64 |
| 9 | battery | 3448 non-null | float64 |
| 10 | weight | 3447 non-null | float64 |
| 11 | release_year | 3454 non-null | int64 |
| 12 | days_used | 3454 non-null | int64 |
| 13 | normalized_used_price | 3454 non-null | float64 |
| 14 | normalized_new_price | 3454 non-null | float64 |

dtypes: float64(9), int64(2), object(4)

memory usage: 404.9+ KB



Checking for missing values

```
In [20]: data.isnull().sum() ## Complete the code to check duplicate entries in the data
```



```
Out[20]: brand_name      0
os                      0
screen_size            0
4g                     0
5g                     0
main_camera_mp        179
selfie_camera_mp       2
int_memory             4
ram                    4
battery                6
weight                 7
release_year           0
days_used             0
normalized_used_price   0
normalized_new_price    0
dtype: int64
```

EDA Results - Univariate Analysis

Normalized new price




Normalized used price





Observation – In this sampling observation, the used price and the new price seem to be distributed normally (visually), with outliers on both sides

EDA Results - Univariate Analysis – Cont.


Internal Memory



Screen size




Main camera MP




Observation – In this sampling observation, these variables look skewed to the right

EDA Results - Univariate Analysis – Cont.



Selfie camera MP




RAM



Weight




Battery




Observation – In this sampling observation, these variables look skewed to the right as well

EDA Results - Univariate Analysis – Cont.


Days Used




Release Year



4g




5g




Observation – Not surprisingly, these variables show that there is more representation for older devices in this set, with more days used, and more 4g than 5g

EDA Results - Univariate Analysis – Cont.

Brand Name




OS



Observation – These 2 variables are most likely correlated, as the OS and the brand usually go together. Note that there is a small representation of iPhones, which can indicate that these devices are not resold like Android devices, or there isn't enough data


Multivariate Analysis – Heatmap/Correlation




Observation – We can see a stronger correlation between screen size, battery and weight. Also, there is a strong correlation between the new price and the old price

Multivariate Analysis – RAM and Weight across brands

Brand Name and RAM relation




Brand Name and weight relation where battery > 4500



Observation – There is no relation between the RAM and the weight within the same brand


Multivariate Analysis – Screen size and cameras across brands

Brand Name and screen size > 6 * 2.54




Observation – We can see that the similar brands have the stronger front cameras and a bigger screen.
However, different brands have stronger rear camera (the main camera)

Brand Name and front camera > 8 mp




Brand Name and weight rear camera > 16 mp




Multivariate Analysis – Prices over time and in relation to 4/5g

Price of used devices



Price of used 4g and 5g



Observation – As expected, newer (but used) devices cost more, and 5g is more expensive than 4g that is more expensive than non-4g

Data Preprocessing – Missing Value Imputation

initial missing data

```
brand_name      0
os              0
screen_size     0
4g              0
5g              0
main_camera_mp 179
selfie_camera_mp 2
int_memory      4
ram             4
battery         6
weight          7
release_year    0
days_used      0
normalized_used_price 0
normalized_new_price 0
dtype: int64
```

After adding the median based on release year and brand

```
brand_name      0
os              0
screen_size     0
4g              0
5g              0
main_camera_mp 179
selfie_camera_mp 2
int_memory      0
ram             0
battery         6
weight          7
release_year    0
days_used      0
normalized_used_price 0
normalized_new_price 0
dtype: int64
```

After adding the median based brand name only

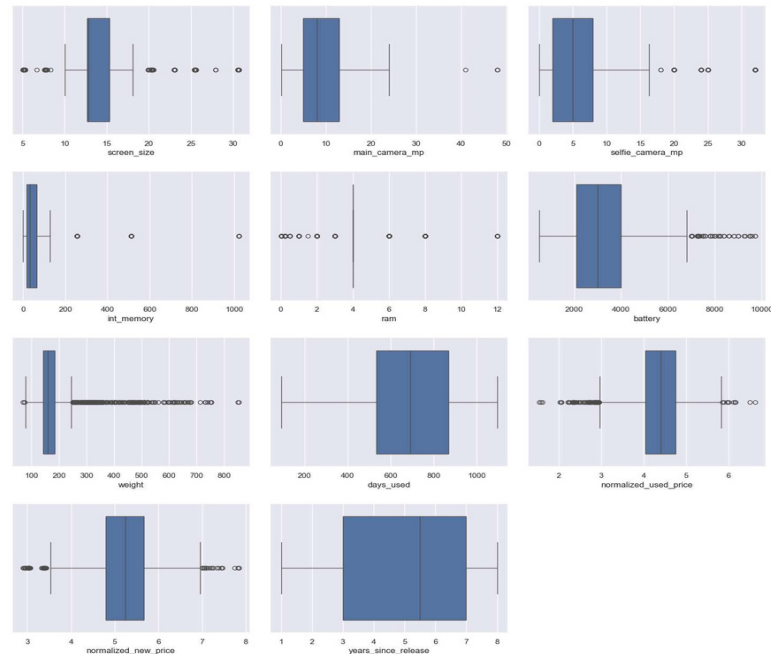
```
brand_name      0
os              0
screen_size     0
4g              0
5g              0
main_camera_mp 10
selfie_camera_mp 0
int_memory      0
ram             0
battery         0
weight          0
release_year    0
days_used      0
normalized_used_price 0
normalized_new_price 0
dtype: int64
```

Final State – After adding the median of the entire column of main_camera_mp

```
brand_name      0
os              0
screen_size     0
4g              0
5g              0
main_camera_mp 0
selfie_camera_mp 0
int_memory      0
ram             0
battery         0
weight          0
release_year    0
days_used      0
normalized_used_price 0
normalized_new_price 0
dtype: int64
```



Data Preprocessing – Feature Engineering and Outlier Check

- New variable was added, to simplify the computation of years.
- Let's create a new column `years_since_release` from the `release_year` column.
- We will consider the year of data collection, 2021, as the baseline.
- We will drop the `release_year` column.



Outliers Check - Most of the variables have outliers, however we will move on with considering them as part of the model

Model Performance Summary – Process Stages



Model Performance Summary – Initial Run Performance

Training Performance

| | RMSE | MAE | R-squared | Adj. R-squared | MAPE |
|---|----------|----------|-----------|----------------|----------|
| 0 | 0.229884 | 0.180326 | 0.844886 | 0.841675 | 0.043268 |

Test Performance

| | RMSE | MAE | R-squared | Adj. R-squared | MAPE |
|---|----------|----------|-----------|----------------|----------|
| 0 | 0.238358 | 0.184749 | 0.842479 | 0.834659 | 0.045017 |

Observation – The results show a strong performance and ability to predict the independent variable (normalized price of used devices).

Now we need to check the assumptions of linear regression.

Model Performance Summary – Checking Linear Regression assumptions

- **NO MULTICOLLINEARITY**

The independent variables must not be highly correlated with one another, as this can lead to unstable and unreliable coefficient estimates.

- **LINEARITY**

The relationship between the independent variables and the dependent variable must be linear.

- **INDEPENDENCE**

The observations in the dataset must be independent of one another, meaning they are not influenced by any other observations in the dataset.

- **NORMALITY**

The residuals must be normally distributed with a mean of zero.

- **HOMOSCEDASTICITY**

The variance of the residuals (the difference between the observed values and the predicted values) must be constant across all levels of the independent

Model Performance Summary – No Multicollinearity

| Index | Feature | VIF |
|-------|-----------------------|------------|
| 0 | const | 227.744081 |
| 1 | screen_size | 7.677290 |
| 2 | main_camera_mp | 2.285051 |
| 3 | selfie_camera_mp | 2.812473 |
| 4 | int_memory | 1.364152 |
| 5 | ram | 2.282352 |
| 6 | battery | 4.081780 |
| 7 | weight | 6.396749 |
| 8 | days_used | 2.660269 |
| 9 | normalized_new_price | 3.119430 |
| 10 | years_since_release | 4.899007 |
| 11 | brand_name_Alcatel | 3.405693 |
| 12 | brand_name_Apple | 13.057668 |
| 13 | brand_name_Asus | 3.332038 |
| 14 | brand_name_BlackBerry | 1.632378 |
| 15 | brand_name_Celkon | 1.774721 |
| 16 | brand_name_Coolpad | 1.468006 |
| 17 | brand_name_Gionee | 1.951272 |

| Index | Feature | VIF |
|-------|----------------------|----------|
| 18 | brand_name_Google | 1.321778 |
| 19 | brand_name_HTC | 3.410361 |
| 20 | brand_name_Honor | 3.340687 |
| 21 | brand_name_Huawei | 5.983852 |
| 22 | brand_name_Infinix | 1.283955 |
| 23 | brand_name_Karbonn | 1.573702 |
| 24 | brand_name_LG | 4.849832 |
| 25 | brand_name_Lava | 1.711360 |
| 26 | brand_name_Lenovo | 4.558941 |
| 27 | brand_name_Meizu | 2.179607 |
| 28 | brand_name_Micromax | 3.363521 |
| 29 | brand_name_Microsoft | 1.869751 |
| 30 | brand_name_Motorola | 3.274558 |
| 31 | brand_name_Nokia | 3.479849 |
| 32 | brand_name_OnePlus | 1.437034 |
| 33 | brand_name_Oppo | 3.971194 |
| 34 | brand_name_Others | 9.711034 |
| 35 | brand_name_Panasonic | 2.105703 |
| 36 | brand_name_Realme | 1.946812 |
| 37 | brand_name_Samsung | 7.539866 |
| 38 | brand_name_Sony | 2.943161 |
| 39 | brand_name_Spice | 1.688863 |
| 40 | brand_name_Vivo | 3.651437 |
| 41 | brand_name_XOLO | 2.138070 |

| Index | Feature | VIF |
|-------|-------------------|-----------|
| 42 | brand_name_Xiaomi | 3.719689 |
| 43 | brand_name_ZTE | 3.797581 |
| 44 | os_Others | 1.859863 |
| 45 | os_Windows | 1.596034 |
| 46 | os_iOS | 11.784684 |
| 47 | 4g_yes | 2.467681 |
| 48 | 5g_yes | 1.813900 |

Observation – The highlighted variables have $VIF > 5$, therefore we will try to remove them one by one and see if it will clear the correlated variables and what will be the impact on the model performance

Model Performance Summary – No Multicollinearity – Cont.

| Column | Adj. R-squared after dropping col | RMSE after dropping col |
|--------------------|-----------------------------------|-------------------------|
| brand_name_Apple | 0.841809 | 0.232201 |
| brand_name_Huawei | 0.841808 | 0.232201 |
| brand_name_Others | 0.841806 | 0.232203 |
| os_iOS | 0.841795 | 0.232211 |
| brand_name_Samsung | 0.841774 | 0.232227 |
| screen_size | 0.838381 | 0.234703 |
| weight | 0.838071 | 0.234928 |

Results – The highlighted variables were removed from the model and the model performance remained high

Model Performance Summary – Dropping high p-value variables

We will drop the predictor variables having a p-value greater than 0.05 as they do not significantly impact the target variable.

Since sometimes p-values change after dropping a variable. So, we'll not drop all variables at once.

Instead, we will do the following:

- Build a model, check the p-values of the variables, and drop the column with the highest p-value.
- Create a new model without the dropped feature, check the results

Model Performance Summary – Second Run Results

Results – The second iteration of the model, without the high VIF and the high P-values variables, show strong results, so we will continue working with it

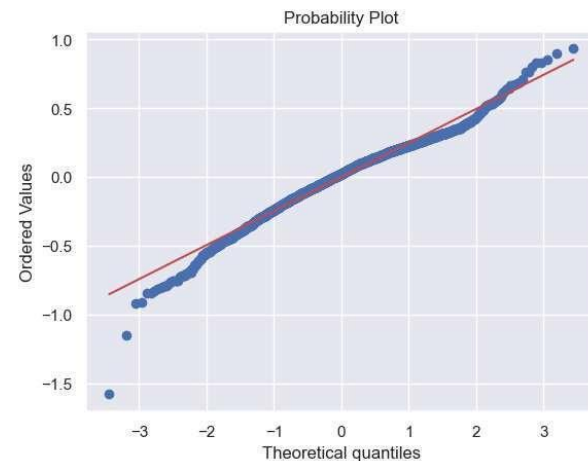


Train performance:

| RMSE | MAE | R-squared | Adj. R-squared | MAPE |
|----------|----------|-----------|----------------|----------|
| 0.249705 | 0.194005 | 0.816984 | 0.81584 | 0.046678 |

Test performance:

| RMSE | MAE | R-squared | Adj. R-squared | MAPE |
|----------|----------|-----------|----------------|----------|
| 0.253286 | 0.196185 | 0.822131 | 0.819518 | 0.047941 |

Model Performance Summary – Linearity and Independence



Observation – By plotting, we see random relation between the fitted and the residuals and a visible normal distribution of the residuals. Even the Q-Q comparison, shows close enough to normal.

Model Performance Summary – Linearity and Independence – Cont.

Last test will be the Shapiro Test:

ShapiroResult(statistic=0.9790064070477261, **pvalue=2.080160689360709e-18**)

Since $p\text{-value} < 0.05$, the residuals are NOT normal as per Shapiro test. Strictly speaking - the residuals are not normal. However, as an approximation, we might be willing to accept this distribution as close to being normal

Model Performance Summary – Homoscedasticity

We will test for homoscedasticity by using the 'goldfeldquandt' test:

Result:

```
('F statistic', 1.0005049517053914), ('p-value', 0.4965187310899025)
```

Since we got a p-value greater than 0.05, we can say that the residuals are homoscedastic. (The null hypothesis was that they are homoscedastic, and we fail to reject it)

Model Performance Summary – Final Model Performance

Training Set Performance:

| RMSE | MAE | R-squared | Adj. R-squared | MAPE |
|----------|----------|-----------|----------------|----------|
| 0.249705 | 0.194005 | 0.816984 | 0.81584 | 0.046678 |

Test Set Performance:

| RMSE | MAE | R-squared | Adj. R-squared | MAPE |
|----------|----------|-----------|----------------|----------|
| 0.253286 | 0.196185 | 0.822131 | 0.819518 | 0.047941 |

Summary – After adjusting the training and test sets to meet the assumptions, the model keeps its strong performance and shows R-squared of 82% which satisfies our expectations.

APPENDIX

Full Data Set Dictionary

The 2021 data contains the different attributes of used/refurbished phones and tablets.

| Name | Description |
|-----------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| brand_name | Name of manufacturing brand |
| os | OS on which the device runs |
| screen_size | Size of the screen in cm |
| 4g | Whether 4G is available or not |
| 5g | Whether 5G is available or not |
| main_camera_mp | Resolution of the rear camera in megapixels |
| selfie_camera_mp | Resolution of the front camera in megapixels |
| int_memory | Amount of internal memory (ROM) in GB |
| ram | Amount of RAM in GB |
| battery | Energy capacity of the device battery in mAh |
| weight | Weight of the device in grams |
| release_year | Year when the device model was released |
| days_used | Year when the device model was released |
| normalized_new_price | Normalized price of a new device of the same model in euros |
| normalized_used_price | This will be the independent variable <small>Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.</small> |

Model Performance Summary – Initial Run Full Results

OLS Regression Results

```

=====
Dep. Variable:    normalized_used_price    R-squared:                0.845
Model:            OLS                    Adj. R-squared:           0.842
Method:           Least Squares          F-statistic:             268.7
Date:             Sun, 25 Aug 2024        Prob (F-statistic):       0.00
Time:             12:21:20                Log-Likelihood:          123.85
No. Observations: 2417                   AIC:                     -149.7
Df Residuals:     2368                   BIC:                     134.0
Df Model:         48
Covariance Type:  nonrobust
=====

```

| | coef | std err | t | P> t | [0.025 | 0.975] |
|-----------------------|------------|----------|--------|-------|-----------|-----------|
| const | 1.3156 | 0.071 | 18.454 | 0.000 | 1.176 | 1.455 |
| screen_size | 0.0244 | 0.003 | 7.163 | 0.000 | 0.018 | 0.031 |
| main_camera_mp | 0.0208 | 0.002 | 13.848 | 0.000 | 0.018 | 0.024 |
| selfie_camera_mp | 0.0135 | 0.001 | 11.997 | 0.000 | 0.011 | 0.016 |
| int_memory | 0.0001 | 6.97e-05 | 1.651 | 0.099 | -2.16e-05 | 0.000 |
| ram | 0.0230 | 0.005 | 4.451 | 0.000 | 0.013 | 0.033 |
| battery | -1.689e-05 | 7.27e-06 | -2.321 | 0.020 | -3.12e-05 | -2.62e-06 |
| weight | 0.0010 | 0.000 | 7.480 | 0.000 | 0.001 | 0.001 |
| days_used | 4.216e-05 | 3.09e-05 | 1.366 | 0.172 | -1.84e-05 | 0.000 |
| normalized_new_price | 0.4311 | 0.012 | 35.147 | 0.000 | 0.407 | 0.455 |
| years_since_release | -0.0237 | 0.005 | -5.193 | 0.000 | -0.033 | -0.015 |
| brand_name_Alcatel | 0.0154 | 0.048 | 0.323 | 0.747 | -0.078 | 0.109 |
| brand_name_Apple | -0.0038 | 0.147 | -0.026 | 0.980 | -0.292 | 0.285 |
| brand_name_Asus | 0.0151 | 0.048 | 0.314 | 0.753 | -0.079 | 0.109 |
| brand_name_BlackBerry | -0.0300 | 0.070 | -0.427 | 0.669 | -0.168 | 0.108 |
| brand_name_Celkon | -0.0468 | 0.066 | -0.707 | 0.480 | -0.177 | 0.083 |
| brand_name_Coolpad | 0.0209 | 0.073 | 0.287 | 0.774 | -0.122 | 0.164 |
| brand_name_Gionee | 0.0448 | 0.058 | 0.775 | 0.438 | -0.068 | 0.158 |
| brand_name_Google | -0.0326 | 0.085 | -0.385 | 0.700 | -0.199 | 0.133 |
| brand_name_HTC | -0.0130 | 0.048 | -0.270 | 0.787 | -0.108 | 0.081 |
| brand_name_Honor | 0.0317 | 0.049 | 0.644 | 0.520 | -0.065 | 0.128 |
| brand_name_Huawei | -0.0020 | 0.044 | -0.046 | 0.964 | -0.089 | 0.085 |

| | coef | std err | t | P> t | [0.025 | 0.975] |
|----------------------|---------|---------|--------|-------|--------|--------|
| brand_name_Infinix | 0.1633 | 0.093 | 1.752 | 0.080 | -0.019 | 0.346 |
| brand_name_Karbons | 0.0943 | 0.067 | 1.405 | 0.160 | -0.037 | 0.226 |
| brand_name_LG | -0.0132 | 0.045 | -0.291 | 0.771 | -0.102 | 0.076 |
| brand_name_Lava | 0.0332 | 0.062 | 0.533 | 0.594 | -0.089 | 0.155 |
| brand_name_Lenovo | 0.0454 | 0.045 | 1.004 | 0.316 | -0.043 | 0.134 |
| brand_name_Meizu | -0.0129 | 0.056 | -0.230 | 0.818 | -0.123 | 0.097 |
| brand_name_Micromax | -0.0337 | 0.048 | -0.704 | 0.481 | -0.128 | 0.060 |
| brand_name_Microsoft | 0.0952 | 0.088 | 1.078 | 0.281 | -0.078 | 0.268 |
| brand_name_Motorola | -0.0112 | 0.050 | -0.226 | 0.821 | -0.109 | 0.086 |
| brand_name_Nokia | 0.0719 | 0.052 | 1.387 | 0.166 | -0.030 | 0.174 |
| brand_name_OnePlus | 0.0709 | 0.077 | 0.916 | 0.360 | -0.081 | 0.223 |
| brand_name_Oppo | 0.0124 | 0.048 | 0.261 | 0.794 | -0.081 | 0.106 |
| brand_name_Others | -0.0080 | 0.042 | -0.190 | 0.849 | -0.091 | 0.075 |
| brand_name_Panasonic | 0.0563 | 0.056 | 1.008 | 0.314 | -0.053 | 0.166 |
| brand_name_Realme | 0.0319 | 0.062 | 0.518 | 0.605 | -0.089 | 0.153 |
| brand_name_Samsung | -0.0313 | 0.043 | -0.725 | 0.469 | -0.116 | 0.053 |
| brand_name_Sony | -0.0616 | 0.050 | -1.220 | 0.223 | -0.161 | 0.037 |
| brand_name_Spice | -0.0147 | 0.063 | -0.233 | 0.816 | -0.139 | 0.109 |
| brand_name_Vivo | -0.0154 | 0.048 | -0.318 | 0.750 | -0.110 | 0.080 |
| brand_name_XOLO | 0.0152 | 0.055 | 0.277 | 0.782 | -0.092 | 0.123 |
| brand_name_Xiaomi | 0.0069 | 0.048 | 1.006 | 0.071 | -0.007 | 0.181 |
| brand_name_ZTE | -0.0057 | 0.047 | -0.121 | 0.904 | -0.099 | 0.087 |
| os_Others | -0.0510 | 0.033 | -1.555 | 0.120 | -0.115 | 0.013 |
| os_Windows | -0.0207 | 0.045 | -0.459 | 0.646 | -0.109 | 0.068 |
| os_iOS | -0.0663 | 0.146 | -0.453 | 0.651 | -0.354 | 0.221 |
| 4g_yes | 0.0528 | 0.016 | 3.326 | 0.001 | 0.022 | 0.084 |
| 5g_yes | -0.0714 | 0.031 | -2.268 | 0.023 | -0.133 | -0.010 |

Model Performance Summary – Final Results

OLS Regression Results

| | | | | | | |
|----------------------|-----------------------|---------------------|---------|-------|----------|--------|
| ===== | | | | | | |
| Dep. Variable: | normalized_used_price | R-squared: | 0.817 | | | |
| Model: | OLS | Adj. R-squared: | 0.816 | | | |
| Method: | Least Squares | F-statistic: | 765.9 | | | |
| Date: | Sun, 25 Aug 2024 | Prob (F-statistic): | 0.00 | | | |
| Time: | 13:43:32 | Log-Likelihood: | -76.051 | | | |
| No. Observations: | 2417 | AIC: | 182.1 | | | |
| Df Residuals: | 2402 | BIC: | 269.0 | | | |
| Df Model: | 14 | | | | | |
| Covariance Type: | nonrobust | | | | | |
| ===== | | | | | | |
| | coef | std err | t | P> t | [0.025 | 0.975] |
| ----- | | | | | | |
| const | 1.3946 | 0.051 | 27.469 | 0.000 | 1.295 | 1.494 |
| main_camera_mp | 0.0147 | 0.001 | 10.594 | 0.000 | 0.012 | 0.017 |
| selfie_camera_mp | 0.0117 | 0.001 | 10.374 | 0.000 | 0.009 | 0.014 |
| ram | 0.0187 | 0.005 | 3.479 | 0.001 | 0.008 | 0.029 |
| battery | 9.434e-05 | 4.99e-06 | 18.911 | 0.000 | 8.46e-05 | 0.000 |
| normalized_new_price | 0.4756 | 0.011 | 42.607 | 0.000 | 0.454 | 0.498 |
| years_since_release | -0.0215 | 0.003 | -6.208 | 0.000 | -0.028 | -0.015 |
| brand_name_Alcatel | 0.0599 | 0.028 | 2.109 | 0.035 | 0.004 | 0.116 |
| brand_name_Karbonn | 0.1408 | 0.058 | 2.414 | 0.016 | 0.026 | 0.255 |
| brand_name_Lenovo | 0.0645 | 0.023 | 2.783 | 0.005 | 0.019 | 0.110 |
| brand_name_Microsoft | 0.1394 | 0.070 | 1.996 | 0.046 | 0.002 | 0.276 |
| brand_name_Nokia | 0.0715 | 0.033 | 2.157 | 0.031 | 0.007 | 0.136 |
| brand_name_Xiaomi | 0.0736 | 0.028 | 2.676 | 0.007 | 0.020 | 0.128 |
| os_Others | -0.1455 | 0.030 | -4.805 | 0.000 | -0.205 | -0.086 |
| 5g_yes | -0.0763 | 0.033 | -2.348 | 0.019 | -0.140 | -0.013 |



Happy Learning !

